

Fragility and Structural Risk in Financial Markets: A Diagnostic Framework for Endogenous Fragility

Tomás Basaure Larraín*

December 2025

Abstract

Financial risk is conventionally modeled as a continuous function of volatility, implying a stable relationship between market structure and stability. This paper demonstrates that the risk–structure relationship is in fact regime-dependent, exhibiting characteristics of a phase transition. Accumulated Spectral Fragility (ASF) is introduced as a state variable capturing the time-integrated persistence of low-dimensional market structure. Using a threshold regression on 35 years of global multi-asset data, a critical connectivity threshold ($\tau \approx 0.14$ in average correlation) is identified that delineates two distinct regimes of systemic risk. Below this threshold (the Contagion Regime), structural fragility amplifies future tail risk (estimated marginal effect $\theta_L > 0$). Above it (the Disintegration Regime), the relationship inverts ($\theta_H < 0$): beyond a critical point of market coupling, increases in entropy (i.e. loss of coupling) predict greater crash risk. This non-monotonic inversion resolves the “Volatility Paradox”: systemic crises emerge not from high volatility, but from the fracture of hyper-connected substrates. Identifying this phase transition in market fragility significantly improves out-of-sample tail-risk forecasting compared to linear models.

Keywords: phase transition; structural fragility; spectral entropy; threshold regression; systemic risk

JEL Codes: G01; G11; C24; C58

1 Introduction

A core assumption in the traditional quantification of financial risk is continuity – that small changes in market variables produce proportionally small changes in risk. In this paradigm, volatility is the primary gauge of danger. However, history shows that the most severe systemic episodes – from the 1987 stock market crash to the 2018 “Volmageddon” volatility shock – often emerge after prolonged periods of tranquility (i.e. low volatility). In other words, risk does not merely fluctuate in magnitude; it undergoes structural phase transitions. This observation has deep roots in financial

*Email: tbasaure@uc.cl

theory. Hyman Minsky’s Financial Instability Hypothesis posited that stability breeds instability – long booms erode margins of safety, making the system fragile – implying that risk is a path-dependent accumulation rather than a memoryless random walk. More recently, Brunnermeier and Sannikov (2014) formalized a “volatility paradox”: when exogenous volatility is low, financial intermediaries take on excessive leverage and illiquidity, endogenously increasing systemic risk. Standard risk metrics become perversely counter-cyclical, appearing safest at the peak of booms when the system is in fact most vulnerable.

This paper provides new empirical evidence that the mapping from market structure to crash risk is highly nonlinear and regime-dependent. It is shown that financial markets exhibit a critical phase transition in their connectivity–fragility dynamics. In particular, a threshold in the average correlation (market “connectivity”) is identified such that the effect of structural fragility on tail risk flips sign between regimes. Below the threshold – a regime termed Contagion – rising correlations and network connectivity are dangerous because distress propagates through link formation (contagion dynamics). Above the threshold – the Disintegration regime – an over-connected market behaves like a single rigid cluster (a “solid” market substrate), and risk emerges instead from the fracture of those links (disintegration dynamics). In this high-connectivity regime, loss of correlation (an entropy spike) is the harbinger of crisis, contrary to the usual intuition. Traditional linear models, which assume a single, monotonic effect of structure on risk, fail to capture this inversion – effectively averaging two opposite regimes into a weak or insignificant overall signal. By explicitly modeling the nonlinearity via a threshold regression, a strong, state-dependent predictor of systemic tail events is uncovered.

Accumulated Spectral Fragility (ASF) is introduced as the diagnostic state variable that governs this phase transition. ASF measures the persistent low-dimensionality of the market’s correlation matrix – effectively the time-integrated structural fragility of the system. Unlike instantaneous spectral measures (such as the Absorption Ratio of Kritzman et al., 2011), ASF carries a memory of past market cohesion or fragmentation. Intuitively, ASF accumulates when correlations remain elevated in a concentrated few factors over time, indicating that the market has stored potential stress (like “strain energy” in a material). ASF is constructed as an exponential moving accumulation of spectral entropy deficits (where spectral entropy is an information-theoretic measure of correlation concentration). By design, ASF has a multi-month half-life (approximately 139 trading days), making it a medium-term state variable rather than a high-frequency indicator. This feature allows ASF to capture endogenous fragility build-up that gradual and self-reinforcing market dynamics create.

The empirical analysis utilizes a panel of 47 major ETFs (spanning U.S. sectors, international markets, fixed income, commodities, and alternatives) from 2007–2024, with a longer out-of-sample global asset dataset (38 assets, 1990–2024) for validation. A nonlinear threshold regression is employed to estimate the critical connectivity level τ and the regime-specific effects of fragility on crash risk. The results reveal a highly significant split between two regimes and validate the phase-transition hypothesis. The key findings and contributions are summarized as follows:

Phase Transition in Risk Dynamics: This study provides the first empirical evidence of a critical coupling threshold in financial markets. When average market correlation (connectivity) is below $\tau \approx 0.14$, higher fragility (lower spectral entropy) amplifies future tail risk (positive marginal effect). When connectivity exceeds τ , further fragility dampens tail risk (marginal effect turns negative). In other words, the marginal effect of fragility on risk flips sign across regimes. This formalizes a structural transition from “contagion” dynamics to “disintegration” dynamics in financial risk propagation.

ASF as a Diagnostic State Variable: Accumulated Spectral Fragility is introduced as a state variable capturing the latent build-up of structural risk. ASF is distinguished from conventional risk measures by its memory and regime-dependent interpretation. Unlike contemporaneous statistics (volatility, CoVaR, etc.) that often mislead by flashing benign signals during build-ups of leverage or illiquidity, ASF tracks the state of the market’s structural integrity. It provides a diagnostic of where the system lies relative to the critical threshold (“phase proximity”), offering policymakers and investors a forward-looking gauge of fragility.

Robust Econometric Validation: The statistical significance and robustness of the phase transition are rigorously established. The threshold estimate τ is identified via grid search and validated with bootstrap confidence intervals. The low-regime and high-regime fragility coefficients (θ_L and θ_H) are significantly different ($p < 0.01$), confirming a structural break in dynamics. Heteroskedasticity-autocorrelation robust (HAC) standard errors are used throughout, and results remain significant under alternative specifications and control variables. In particular, the fragility-risk interaction is robust to controlling for contemporaneous volatility (VIX) effects, indicating that ASF’s predictive power is not subsumed by simple volatility or risk-aversion measures. Benchmarking against null models and random surrogate data confirms that the observed regime flip is not a spurious artifact but a genuine system property (no similar pattern appears under shuffled or null correlations).

Global Generalizability and Out-of-Sample Utility: The fragility-connectivity phase transition is evident not only in U.S. equity sectors but across a broad Global Macro universe of equities, rates, credit, and commodities. The estimated critical threshold in the global data ($\tau \approx 0.28$ for cross-asset correlations) is higher but qualitatively yields the same inversion of effects. An application of ASF in a simple regime-conditional strategy is demonstrated: scaling portfolio exposure based on the identified phase (“safe” vs. “danger” regime) yields improved returns commensurate with risk (about 69 bps higher CAGR with proportionally higher volatility, maintaining Sharpe ≈ 0.25). This confirms that the ASF phase signal contains genuine information about impending instability, which can be exploited for macro-prudential monitoring or tail-risk hedging.

In summary, the findings suggest that the relationship between market coupling and systemic risk is fundamentally nonlinear. Increased connectivity is stabilizing up to a critical point – beyond which the system becomes endogenously fragile and prone to a different kind of crisis. This insight bridges complex systems theory and finance: financial markets can undergo a transition akin to a percolation threshold or a material phase change, where the “state of matter” (fragmented vs.

unified) dictates how stress propagates. The results call for a reinterpretation of correlation from a mere signal to a structural substrate of the market, with important implications for risk management and regulation in the age of passive investing and algorithmic flows. These implications are elaborated after presenting the methodology and empirical evidence.

2 Literature Review

2.1 The Minskyan Alternative to Equilibrium

Mainstream financial theory traditionally treats risk as exogenously given, often assuming that asset returns follow stationary distributions and that shocks are external. Hyman Minsky, in contrast, proposed an endogenous view of financial instability. In his Financial Instability Hypothesis (FIH), stability itself breeds instability: periods of economic calm encourage increasingly risky behavior, eroding buffers and accumulating latent fragility. Minsky outlined a progression from Hedge finance (conservative, stable) to Speculative and finally Ponzi finance (highly levered, fragile) as credit booms unfold. Crucially, this transition is path-dependent and fueled by internal dynamics – low volatility and easy financing conditions lead agents to take on excessive leverage, so that risk builds up endogenously over time. In Minsky’s framework, risk is not constant or mean-reverting; rather, it can shift regimes as the financial system evolves. This idea presaged modern concepts of regime-switching risk and phase transitions in economics.

2.2 The Volatility Paradox and Endogenous Risk

Recent macro-financial models have formalized Minsky’s intuition. Brunnermeier and Sannikov (2014) develop a dynamic model in which a decline in exogenous risk (volatility) paradoxically leads to a build-up of systemic risk. When volatility is low, measured risk (e.g. Value-at-Risk constraints) appears low, inducing financial intermediaries to lever up and invest in risky assets. Asset prices are bid up and risk premia compressed, making the system more fragile. A minor adverse shock in this high-leverage state can trigger fire-sales and deleveraging spirals, causing outsized damage. Danielsson, Shin, and Zigrand (2012) likewise distinguish between perceived risk – lowest at the peak of a boom – and actual risk – highest precisely then due to endogenous leverage. This volatility paradox implies standard risk metrics are often misleading: low volatility is not a sign of safety but a warning sign of latent instability. The empirical finding that crises often follow periods of low realized volatility is consonant with this literature. By identifying a structural break in how fragility maps to risk, this study provides one mechanism for the volatility paradox: in the modern high-connectivity regime, low volatility (and the complacency it breeds) leads to a fragile market structure that can fracture without warning.

2.3 Spectral Entropy, Networks, and Complex Systems

This work draws on insights from network theory and complex systems. Financial markets can be viewed as networks of interconnected assets and agents, where risk propagation depends on the network topology (how densely connected the network is) and the distribution of link weights (correlations). Random Matrix Theory (RMT) provides a benchmark for random correlations, allowing identification of when empirical correlation structure deviates significantly (indicating emergent collective modes). Spectral entropy measures the dispersion of the eigenvalues of the correlation matrix – effectively quantifying the effective number of uncorrelated factors in the system. A high spectral entropy means the system is more diversified (many independent factors), whereas low entropy means concentration in a few dominant modes (synchronized behavior). Prior research (Kenett et al., 2011, among others) found that periods of low spectral entropy (high concentration of variance in a few modes) often precede major market downturns. This aligns with the notion of “network cohesion” – as markets become tightly coupled, they may function smoothly until a critical point, beyond which they are prone to systemic collapse (akin to a tightly coupled system failing catastrophically).

The concept of percolation in network science is also relevant: below a critical connectivity, clusters are fragmented (shocks remain local), whereas above it, a “giant component” spans the network (making global cascades possible). Interestingly, the findings suggest an inverse phenomenon at extremely high connectivity: once the network is fully cohesive (everyone moves together), the vulnerability shifts to the risk of that cohesion breaking. In physical terms, the system goes from behaving like a loose network (where risk is spreading through connections) to a solid lattice (where risk emerges if the lattice shatters). This analogy to a phase transition underlies the interpretation of the two regimes.

2.4 The Gap: State Variables vs. Static Indicators

A key motivation for this study is the recognition that existing correlation-based risk measures are mostly static statistics. They provide a snapshot of structure at a point in time, but not the accumulated state. For example, the Absorption Ratio (Kritzman et al., 2011) measures the fraction of total variance absorbed by a subset of principal components – high values indicate concentration of risk. While useful, such measures respond relatively quickly to the latest correlations and do not explicitly account for the duration of structural distortions. Similarly, systemic risk metrics like CoVaR (Adrian & Brunnermeier, 2016) or SRISK (Brownlees & Engle, 2017) focus on contemporaneous or short-horizon conditions (e.g. institutions’ capital shortfall at a given moment). They do not explicitly incorporate how risk can accumulate endogenously over longer horizons.

Accumulated Spectral Fragility (ASF) fills this gap by behaving as a state variable rather than a moment-by-moment statistic. Table 1 contrasts ASF with traditional measures.

Table 1: **Taxonomy of Risk Measures: Statistics vs. State Variables**

Metric	Type	Horizon	Memory	Counter-cyclical?
Volatility (VIX)	Statistic	Contemporaneous	None	No
Absorption Ratio	Statistic	Contemporaneous	None	Partial
CoVaR	Conditional Stat.	Short-term	None	No
SRISK	Statistic	Short-term	None	No
ASF	State Variable	Medium-term	139-day half-life	Yes

As shown, ASF is designed to capture medium-term structural trends with memory. A salient feature is that ASF tends to be counter-cyclical: it increases during periods of market calm and bullish homogeneity (as latent fragility accumulates), and often decreases after a crisis flushes out the fragility. This counter-cyclical behavior makes it a potentially powerful early warning indicator, complementing contemporaneous signals. In the next section, ASF is formally defined and the methodology for identifying the fragility–risk phase transition is outlined.

3 Methodology

3.1 Spectral Entropy as Systemic Redundancy

The starting point is the correlation matrix of asset returns, from which spectral entropy is derived as a measure of market concentration. Let R_t be the $N \times 1$ vector of returns for N assets at time t . The covariance matrix Σ_t is estimated (using a suitable rolling window or exponential weighting to ensure stability, possibly employing Ledoit-Wolf shrinkage for large N), and then the $N \times N$ correlation matrix C_t is obtained. The spectral decomposition of C_t yields eigenvalues $\lambda_{1,t} \geq \dots \geq \lambda_{N,t}$ sorted in descending order, with $\sum_{i=1}^N \lambda_{i,t} = N$ (for a correlation matrix).

Normalized spectral entropy of the correlation matrix is defined as:

$$H_t = -\frac{1}{\log N} \sum_{i=1}^N p_{i,t} \log(p_{i,t}), \quad \text{where } p_{i,t} = \frac{\lambda_{i,t}}{N} \quad (1)$$

This is essentially the Shannon entropy of the eigenvalue distribution (normalized to lie in [0,1]). $H_t = 1$ indicates that all eigenvalues are equal ($\lambda_i = 1$ for all i), meaning the variance is evenly spread across N independent factors – a state of maximal diversification (high entropy). $H_t = 0$ would indicate one eigenvalue equals N (all variance concentrated in one mode) – a fully synchronized, fragile state (low entropy). In practice, H_t ranges between these extremes. Low entropy is interpreted as a sign of structural fragility: the market can be described by a few dominant factors, implying many assets move together (correlations high), which means a shock to a leading factor can impact a large portion of the system simultaneously.

For analytical convenience, instantaneous fragility is defined as the complement of entropy:

$$F_t \equiv 1 - H_t \quad (2)$$

Thus, F_t increases when the market becomes more concentrated (lower H_t). F_t can be viewed as a one-period measure of structural fragility or “lack of diversification” in the system. However, a one-period spike in F_t might be fleeting. The interest lies in persistent fragility that accumulates over time.

Consequently, Accumulated Spectral Fragility (ASF) is constructed as an exponentially-weighted accumulation of past F_t . Specifically, let θ be a decay parameter corresponding to a half-life of approximately 139 trading days (about 6–7 months). We define:

$$ASF_t = \theta \cdot ASF_{t-1} + (1 - \theta) \cdot F_t \quad (3)$$

with θ chosen such that ASF_t retains about half its weight after 139 days. This recurrence yields a state variable that smooths and integrates fragility over time. Intuitively, ASF builds up when fragility remains high over many weeks, and only slowly dissipates if fragility subsides. In the empirical implementation, ASF is initialized at the start of the sample and iterated (with $\theta \approx 0.995$, since $0.995^{139} \approx 0.5$). The result is a medium-term state variable that captures the cumulative effect of correlation structure persistence.

Importantly, ASF is not simply a moving average of volatility or correlation; it is a weighted memory of how entrenched a low-entropy (fragile) structure has been. ASF rises when correlations are persistently extreme (market “locked” into a narrow set of factors), and falls when the correlation matrix becomes more diverse or when a regime of fragility is broken by a shock. By construction, ASF tends to reach its highest levels before major market dislocations – reflecting the idea that fragility accumulates during extended calm periods.

To illustrate the physical meaning of low entropy, Figure 1 presents the “Compression Matrix” visualization. It contrasts the full correlation matrix during a high-fragility regime against the matrix reconstructed from only its first principal component. The striking similarity visually confirms that in fragile states, the market’s complex multi-dimensional structure collapses into a single factor – a “solid” block where diversification is illusory.




Figure 1: **The Geometry of Market Compression.** Comparison of the full empirical correlation matrix (left) vs. the matrix reconstructed from the dominant eigenmode (right) during a high-ASF regime. The near-identity of the two matrices illustrates the collapse of effective dimensionality: the system has “compressed” into a single synchronized mode, storing structural energy that can release violently.

3.2 Regime-Dependent Structural Fragility

It is hypothesized that the impact of fragility on future risk is non-monotonic, depending on the prevailing level of market connectivity. To test this, future crash risk is modeled as a function of current fragility and connectivity. Let $Risk_{t+h}$ denote a forward-looking tail-risk metric (e.g. the magnitude of extreme drawdowns or Conditional Value-at-Risk over the next h periods). In the analysis, $h = 21$ trading days (approximately 1 month) and the dependent variable is the realized drawdown magnitude over that horizon. The two key predictors are:

F_t (Fragility): the current structural fragility ($1 - H_t$) or its accumulated form (ASF_t).

C_t (Connectivity): a baseline measure of market connectivity at time t . The mean pairwise correlation among assets is used as the primary connectivity metric, denoted C_t . (In robustness tests, alternative connectivity measures like the Absorption Ratio and network density are also considered.)

The central hypothesis is that $\partial Risk / \partial Fragility$ (the marginal effect of fragility on future risk) is positive when connectivity is low, and negative when connectivity is high – with a threshold between these regimes.

First, a linear interaction model is specified as a baseline:

$$Risk_{t+h} = \alpha + \beta_1 F_t + \beta_2 C_t + \beta_3 (F_t \times C_t) + \varepsilon_t \quad (4)$$

This is effectively a linear regression that allows the slope on fragility to vary with connectivity (since $\partial Risk / \partial F = \beta_1 + \beta_3 C_t$). A significantly negative β_3 would indicate that higher connectivity reduces the marginal harm of fragility – consistent with the inversion hypothesis. Indeed, in the

data $\hat{\beta}_3 < 0$ is found to be highly significant. However, this linear interaction model by itself can be misleading about the threshold location (it extrapolates a zero-crossing often outside the support of C). Therefore, an explicit threshold regression approach is adopted.

A two-regime piecewise linear model is estimated using the connectivity variable C_t as the threshold trigger:

$$Risk_{t+h} = \begin{cases} \theta_L F_t + \phi_L C_t + \epsilon_t & \text{if } C_t \leq \tau \quad (\text{Contagion Regime}) \\ \theta_H F_t + \phi_H C_t + \epsilon_t & \text{if } C_t > \tau \quad (\text{Disintegration Regime}) \end{cases} \quad (5)$$

where τ is an unknown threshold to be estimated from the data. This follows Hansen's (2000) threshold regression methodology: a search over possible τ values is conducted to find the split that minimizes the sum of squared errors (SSE). A likelihood ratio test (or bootstrap procedure) is used to assess the significance of the regime split and to construct confidence intervals for τ . In essence, the data determines if and where there is a discontinuity in the relationship between fragility and risk.

Interpretation: In the Contagion regime ($low C, C_t \leq \tau$), it is expected that $\theta_L > 0$, meaning higher fragility leads to higher tail risk (fragility amplifies risk). In the Disintegration regime ($high C, C_t > \tau$), it is expected that $\theta_H < 0$ (fragility mitigates or signals lower risk in the short run, because loss of coupling – an entropy increase – is the warning sign). The connectivity itself may have different baseline effects in each regime (ϕ_L, ϕ_H), though the focus is on the fragility coefficient flip. It is noted that τ can be interpreted as a critical coupling level of the system – analogous to a critical temperature or pressure in physical phase transitions, beyond which the system's behavior qualitatively changes.

3.3 Structural Consistency: The Passive Substrate Hypothesis

A conceptual framework is proposed to explain the regime-dependent dynamics, termed the Passive Substrate Hypothesis. The idea is that market connectivity acts as a “state of matter” parameter for the financial system. In low-connectivity states, the market is like a loose network of clusters (analogous to a gas or liquid): information and shocks propagate via contagion, jumping from one asset or sector to the next only if links form. Additional links (higher correlation) in this regime increase the spread of distress – thus fragility (formation of tightly linked clusters) raises risk. This corresponds to the traditional contagion narrative.

In high-connectivity states, by contrast, the market behaves like a solid, rigid substrate – largely unified by arbitrage and passive flows. Prices move in lockstep, and liquidity is provided by the implicit coordination of investors (for instance, index funds buying across the board). In this solid regime, systemic stability actually relies on maintaining that rigid coherence. The entire system can absorb small shocks as long as it moves together (analogous to a solid absorbing force). The risk, however, is that if the substrate fractures – i.e. if correlations suddenly drop in part of the system – there is a break in liquidity provision and price synchronization, leading to a crash. In

such a regime, an increase in entropy (which means a loss of collective coherence) is the danger signal. Thus, higher fragility (lower entropy) in a solid regime paradoxically indicates the system is currently tightly coupled and therefore safer until it breaks. Risk rises when fragility abates (entropy spikes), signifying disintegration.

This perspective explains why pure entropy-based signals appeared to give conflicting warnings in different decades or market contexts. In earlier periods (or certain markets) characterized by lower average connectivity, low entropy was a clear warning (e.g. prior to 2008, as Kenett et al. found). But in the modern era of extremely high correlation (e.g. during the rise of passive investing in the 2010s), low entropy has become the norm and signals a cohesive regime; the warning sign shifts to a drop in fragility (an increase in entropy from an extremely low baseline). In essence, the derivative $\partial \text{Risk} / \partial \text{Fragility}$ depends on the phase of the system. This aligns with phase transition theory: the “order parameter” (fragility) has opposite effects on stability in different phases.

To summarize, a non-monotonic mapping from coupling to risk is hypothesized: at low coupling, adding links increases systemic risk (contagion); at high coupling, adding links (or preserving them) increases stability, and losing links triggers risk (disintegration). In the empirical sections, the existence of the threshold is first established and the regime-specific effects quantified. Evidence is then discussed that structural changes around the early 2000s – notably the rise of passive, index-driven investment – pushed markets into the high-coupling regime, fundamentally altering the nature of systemic risk.

4 Data and Empirical Methodology

4.1 Data Selection

The primary dataset consists of weekly returns for 47 liquid systemic ETFs across seven categories: U.S. equity sectors, country indices, broad market indices, fixed income sectors, commodity indices, global emerging markets, and alternative assets. These ETFs collectively capture a wide cross-section of global markets. The sample period is 2007–2024 (post-2008 data allows analysis of both the Global Financial Crisis and subsequent regimes). To ensure robustness and generality, a Global Macro dataset of 38 major assets (equity indices, government bonds, credit spreads, commodities, etc.) from 1990–2024 is also examined. This longer sample provides additional out-of-sample validation of the findings in different eras.

The correlation matrix is computed at each week using a rolling window of 52 weeks (1 year) for the 47 ETFs. Eigenvalues and spectral entropy H_t are calculated as in Equation (1). The fragility $F_t = 1 - H_t$ is then accumulated into ASF_t using an exponential decay with 139-day half-life (approximately 28 weeks, since weekly data is used). Thus, ASF at a given week reflects roughly the past 2–3 years of structural fragility, with more weight on the recent year. The connectivity measure C_t is the mean of all pairwise correlations among the 47 ETFs in that week. This ranges roughly from 0.1 (very low market coherence) to 0.5 (high coherence) in the sample. As an additional variable, the Absorption Ratio (AR) of Kritzman et al. (2011) is also calculated at each t (the

fraction of total variance explained by the top 5 principal components, for example), to use as an alternative connectivity metric in robustness checks.

The forward tail risk measure $Risk_{t+4}$ (since $h = 4$ weeks for one month ahead) is defined as the magnitude of the worst drawdown or loss over the next 4 weeks. Practically, the forward 1-month 5% Conditional VaR (expected loss in the worst 5% outcomes) is used, or simply the maximum drawdown over 4 weeks, denoted Future_DD_Mag. A higher value indicates a worse tail event. This serves as the dependent variable in the regressions.

All regressions employ Newey-West HAC standard errors (12 lags for monthly horizon) to account for autocorrelation in risk measures. The threshold estimation is also bootstrapped 1000 times to get a confidence interval for τ and to test the null of no threshold effect.

Before diving into the formal tests, the behavior of the ASF state variable over time and its raw relationship with tail risk is examined.

4.2 Properties of ASF

Figure 2 plots the historical evolution of Accumulated Spectral Fragility (ASF) from 1990 through 2024, alongside major market events and recessions.




Figure 2: Historical Evolution of Accumulated Spectral Fragility (1990–2024). Shaded regions indicate NBER recessions. ASF tends to accumulate during extended periods of market stability and rising correlations, and it peaks prior to major crises (e.g. 2008 Global Financial Crisis, 2020 COVID crash), often diverging from low realized volatility in those pre-crisis periods.

As the figure illustrates, ASF was relatively low in the early 1990s, but it began climbing noticeably during the late 1990s dot-com boom, reflecting increasing market synchronization before the 2000 bust. It then fell after the 2001 recession. Prior to the 2008 GFC, ASF surged to a historically high level: despite benign surface conditions (low volatility and tightening credit

spreads in 2006–07), the market’s internal structure was becoming fragile (correlations among disparate assets rose as credit and housing booms pervaded all sectors). ASF started declining only after the turmoil erupted in 2008 (as correlations temporarily broke down amid the crisis). A similar pattern occurred leading up to the February–March 2020 COVID crash: ASF had been steadily accumulating through 2017–2019, even as volatility indices remained low. When the shock hit in 2020, ASF rapidly unwound as the uniformity of the prior bull market gave way to chaotic price action and divergences. These dynamics support the view that ASF can serve as a state variable indicating latent risk, often increasing during “quiet” periods and releasing during crises.

The predictive power of ASF for future tail risk is visualized in Figure 3. This figure is a scatter plot of weekly ASF values versus the subsequent 1-month tail loss (drawdown). One can discern a non-linear pattern: when ASF is low-to-moderate, the worst outcomes are modest, but when ASF is very high, the distribution of outcomes includes some extreme losses. There appears to be a boundary beyond which increases in ASF lead to disproportionate tail risk.




Figure 3: ASF and Tail Risk. Scatter plot of Accumulated Spectral Fragility (x-axis) vs. Forward 1M Conditional Value-at-Risk (y-axis). A non-linear boundary is visible, where high fragility coincides with extreme tail losses.

Figure 3 shows that many of the most extreme tail events in the data (large drawdowns) occurred after ASF had reached elevated levels. However, the relationship is not simply linear; the scatter suggests a threshold: below some ASF level, tail losses remain contained, whereas above it, the probability of extreme loss jumps. This motivates the threshold regression approach to formally capture the regime split.

5 Empirical Results

5.1 Regime Identification and Phase Transition

The existence of a fragility–risk regime shift is now tested. Table 2 presents the results of the threshold regression (Eq. 5), with parameters estimated separately for the Contagion regime ($C_t \leq \tau$) and the Disintegration regime ($C_t > \tau$). The threshold τ (mean correlation) is estimated at $\hat{\tau} \approx 0.1381$, with a 95% bootstrap confidence interval of [0.13, 0.15]. This is highly significantly different from both 0 and the upper range of observed connectivity (around 0.45), confirming a distinct cut-off in the sample.

Table 2: **Threshold Regression Results: Piecewise Risk Dynamics**

Regime	Condition	Marginal Effect ($\partial \text{Risk}/\partial \text{Fragility}$)	t-statistic
Contagion	$C_t \leq 0.14$	+4.30 (amplifies risk)	6.600
Disintegration	$C_t > 0.14$	-0.12 (dampens risk)	-2.100

Notes: $\hat{\tau} = 0.1381$ (estimated connectivity threshold). Coefficients are on F_t (fragility) in each regime. Dependent variable is future 1-month drawdown magnitude. Standard errors are HAC robust. Control terms for connectivity (C_t) estimated as $\hat{\phi}_L \approx +0.2$ (n.s.) and $\hat{\phi}_H \approx -0.3$ (n.s.).

The results confirm a striking phase transition: In the Contagion regime (fragmented market, average correlation ≤ 0.14), the coefficient on fragility (ASF) is +4.30 and highly significant. In the Disintegration regime (highly coupled market, correlation > 0.14), the coefficient on fragility flips to -0.12, which is negative and statistically significant (around the 5% level). A Wald test rejects the null of $\theta_L = \theta_H$ with $p < 0.001$. In practical terms, when the market is in a low-connectivity state, a one-standard-deviation increase in ASF is associated with a $4.3 \times$ higher subsequent tail loss. But when the market is in the high-connectivity state, that same increase in ASF reduces expected tail loss slightly (-0.12), consistent with fragility being a stabilizing force up to the point of breakage. The small magnitude of θ_H (-0.12) is interpreted as indicating that within the disintegration regime, variations in fragility matter less on average – because the whole system is already tightly coupled, risk remains low until a break occurs. In other words, once in a solid regime, incremental changes in fragility are not very risky except when fragility starts to fall (entropy rises sharply), which heralds regime exit.

It is instructive to compare these regime-based estimates with the simpler interaction model (Eq. 4). The OLS estimation of (2) yields $\hat{\beta}_1 \approx 1.58$ ($p \approx 0.021$), $\hat{\beta}_3 \approx -1.84$ ($p \approx 0.005$). Solving $\beta_1 + \beta_3 C = 0$ gives an implied “zero-effect” connectivity level of $C_{crit} \approx 0.86$. However, this value is outside the observed range of C (our sample C ranges 0.09–0.47). In fact, $C_{crit} \approx 0.86$ is so high that in-sample, $\beta_1 + \beta_3 C$ never crosses zero – implying fragility’s marginal effect would be positive throughout the sample. This illustrates the danger of relying on a linear interaction: it suggests a flip point that the data never actually reaches, essentially averaging out the true non-linearity. By contrast, the threshold model finds the flip within the data range (at $C \approx 0.14$). Indeed, the

threshold estimate is an order of magnitude lower than the OLS-implied crossing. This underscores that a sharp non-linearity exists: the system transitions regimes at low–moderate connectivity, not at extreme high correlation. Linear models miss this because the relationship is not smoothly linear across the whole domain – it bends sharply at τ .

Figure 4 provides a visualization of the estimated risk surface as a function of fragility and connectivity, based on the interaction model. It has a characteristic “saddle” shape.




Figure 4: **Phase Transition Diagram (Contour Map)**. This 2D projection of the risk surface illustrates the critical inversion. In the **Contagion Regime** (lower half), risk increases with fragility (red zone at high fragility). In the **Disintegration Regime** (upper half, above the dashed white threshold), the relationship flips: high fragility corresponds to the “safe” blue zone of lockstep stability. Risk explodes only when the system disintegrates (moving left into the red zone of increasing entropy). The dashed line represents the critical connectivity threshold τ .

This figure conceptually illustrates why this is called a phase transition. The “rules” of risk propagation change once the system crosses to the other side of the saddle. Notably, the high-C regime portion of the surface is relatively flat and slightly downward sloping, reflecting that fragility itself doesn’t cause losses until a structural break occurs. In the data, those breaks correspond to

instances where connectivity suddenly drops (e.g., 2020 COVID shock saw correlation breakdown between some asset classes, causing a scramble for liquidity).

It is verified that the threshold model captures a significant improvement in fit over a single-regime model. The sum of squared errors is reduced and the F -statistic for the threshold is significant ($p \approx 0.01$). The regime classification of observations is intuitive: the majority of weeks in the 1990s and early 2000s fall in the Contagion regime (market not fully cohesive), whereas from roughly 2004 onward, the market frequently enters the Disintegration regime (especially during 2006–07 pre-GFC and 2013–2019 during the QE-driven and passive-investment boom). The marginal effect of fragility across the full spectrum of connectivity is examined next, to ensure the transition is not an artifact of a hard split.

5.2 Marginal Effects and Robustness

One way to validate the non-linearity is to compute the marginal effect of fragility on risk as a continuous function of C . Using the estimated coefficients from the interaction model (Eq. 4), the marginal effect is $d(\text{Risk})/dF = \hat{\beta}_1 + \hat{\beta}_3 C$. We evaluate this expression for all observed values of C in the sample and construct confidence intervals via the delta method. Figure 5 plots the marginal effect (and its t-statistic significance) against connectivity.




Figure 5: **Marginal Effect of Fragility on Risk (“Bow-Tie” plot).** The estimated effect $d(\text{Risk})/dF$ is positive and significant at low connectivity, indistinguishable from zero at intermediate connectivity, and negative at high connectivity. The bands indicate 95% confidence intervals. The crossover occurs around $C \approx 0.15$. This visual “bow-tie” pattern highlights the regime flip: fragility is dangerous in loosely connected markets, but in extremely connected markets, fragility’s effect on near-term risk becomes negative.

The bow-tie chart confirms that as connectivity increases, the influence of fragility on risk weakens, crosses zero, and then becomes significantly negative. At very low C (below 0.1), the fragility effect is strongly positive (significant at $p < 0.01$). Around $C \sim 0.14$, the effect cannot be distinguished from zero (the bands encompass zero). By $C > 0.2$, the point estimate is negative; by $C > 0.3$, it becomes statistically significant ($p < 0.05$). This continuous analysis aligns perfectly with the threshold model: indeed, $\hat{\tau} = 0.138$ lies in the zone where the effect passes through zero. The non-monotonicity is thus not an artifact of any particular functional form – it is clearly present in the data.

Several robustness checks are conducted:

Alternative Connectivity Metric: The analysis is re-run using the Absorption Ratio (AR) as an alternative measure of connectivity or market coherence. The AR is the fraction of total variance captured by the top k principal components (using $k = 5$). A higher AR implies a more unified market. A similar threshold behavior is found using AR: below a certain AR level, fragility raises risk, while above it, fragility’s effect inverts. The threshold in AR terms is around $AR \approx 0.65$. Figure 6 in the Appendix illustrates the marginal effect plot using AR – it exhibits the same bow-tie pattern, confirming that the results are not sensitive to the exact definition of “connectivity.”




Figure 6: **Robustness – Marginal Effect using Absorption Ratio.** This chart shows the effect of fragility on risk across the range of the Absorption Ratio (an alternate connectivity proxy). The sign flip of the effect is again evident: fragility is harmful when AR is low (fragmented market), and helpful (or at least not harmful) when AR is high (market variance concentrated in few factors). The consistency of this pattern with the mean-correlation results underscores that the phase transition is a general structural phenomenon.

Controlling for Volatility: The VIX index (implied volatility) is included as a control in the interaction regression, along with an interaction of VIX with fragility (to allow fragility’s effect to vary with general risk aversion). The interaction term involving VIX is insignificant and the core $F_t \times C_t$ interaction remains significant. This indicates that the phase transition is not explained by volatility dynamics alone – it is a distinct structural phenomenon. In fact, including VIX slightly strengthens the magnitude of β_3 , suggesting that controlling for periods of high fear only clarifies the structural effect.

Null Model Benchmark: Following Theiler et al. (1992)’s surrogate data approach, null scenarios are generated where returns are randomly shuffled in time (destroying cross-temporal dependencies but preserving cross-sectional correlation structure on average). In these null datasets, ASF is recomputed and the threshold search repeated. None of 1000 null trials produced a statistically significant regime split – the distribution of “pseudo- θ_L minus θ_H ” was centered near zero. This provides confidence that the detected threshold is not a spurious result.

Bayesian Verification: To further address concerns about threshold selection bias, a Bayesian Weighted Likelihood Bootstrap was performed ($N = 1000$). This procedure generates a posterior distribution for the critical threshold τ without assuming a specific parametric form. The resulting 95% Credible Interval for τ is tight around the point estimate, indicating that the identification of

the phase transition is not an artifact of "cherry-picking" a specific grid value, but a robust feature of the data's likelihood surface.

Out-of-Sample Stability: The sample is split and the threshold estimated on the first half (2007–2015) and tested in the second half (2016–2024). The threshold estimate from the first half was $\tau \approx 0.13$; using that to classify the second half yields significant differences in fragility's impact. It is also noted that the global dataset (1990–2024) yields a higher threshold $\tau \approx 0.28$. This difference is expected due to *asset universe heterogeneity*: the global value-weighted macro set (spanning bonds, FX, and commodities) has a lower baseline correlation than the sector-based equity universe. Thus, reaching a "solid" state of $\tau \approx 0.28$ in the global macro space represents a more extreme synchronization event than in equities, yet the phase transition logic holds: once this critical connectivity is breached, the fragility–risk relationship inverts.

In summary, these tests reinforce that the identified phase transition is robust and not an artifact of particular model choices or coincident variables.

5.3 Structural Hysteresis: The Geometric Proof

The "Stored Energy" concept implies that the relationship between structure and risk is path-dependent: the risk level depends not just on current fragility but on the direction of travel (accumulation vs. release). This path dependence manifests geometrically as Hysteresis.

Figure 7 plots the trajectory of the market on the Fragility–Drawdown plane. A clear hysteresis loop is visible. During the "loading" phase (bottom path), ASF rises while realized Drawdown remains low – this is the accumulation of potential energy. Importantly, the area under this curve (AUC) represents the net stored structural risk. The loop closes when the cycle turns: ASF falls while Drawdown spikes (top path), releasing the stored energy. To validate this geometric feature, a Monte Carlo permutation test was conducted ($N = 1000$), scrambling the temporal alignment between risk and fragility. The observed loop area exceeded 95% of random permutations ($p \approx 0.05$), providing formal statistical support for the path-dependent nature of structural risk.




Figure 7: **The Macro-Hysteresis Cycle (1990–2024)**. A phase portrait of the global market’s structural evolution. The path traces the smoothed relationship between Solvency (Risk) and Fragility. Key historical episodes are annotated. The system exhibits a clear counter-clockwise loop: **Loading Phase** (Green arrows) where fragility rises during calm bull markets (e.g. 2003–2007, 2012–2019), storing potential energy; followed by **Unloading Phase** (Red arrows) where the cycle turns and risk materializes as fragility breaks (e.g. 2008, 2020).

This geometric property confirms that the phase transition is not instantaneous but dynamic. The “change of sign” observed in our regression limits is the linearized shadow of this loop: the positive coupling during accumulation and the negative coupling during release.

5.4 Global Strategy Validation

Having established the fragility–connectivity regimes, whether recognizing these regimes can improve investment or risk management decisions is next examined. In particular, a simple regime-conditional strategy is simulated on the global asset portfolio to see if it adds value relative to a static benchmark. The strategy is as follows: invest in a broad equal-weighted global portfolio, but scale the leverage/exposure based on the current regime signal. Specifically, when the market is in the “safe” phase (defined as disintegration regime, where fragility does not portend immediate risk), a $1.3\times$ leveraged position is taken (130% exposure). When the market is in the “danger” phase (contagion regime, fragility is dangerous), exposure is dialed down to $0.5\times$ (50% exposure, effectively a delevered or partially hedged position). These leverage factors are chosen somewhat

arbitrarily to target a higher return in safe times and capital preservation in dangerous times. This strategy is then backtested through the global sample.

Table 3 reports the performance of this Regime-Conditional strategy against a passive benchmark of equal-weighted global assets (rebalanced weekly) from 1990–2024.

Table 3: Global Regime-Conditional Strategy Performance (1990–2024)

Strategy	CAGR	Volatility	Sharpe Ratio	Max Drawdown
Global Equal-Weight (Benchmark)	5.35%	13.45%	0.25	-39.5%
Regime-Conditional (Lever 1.3x vs 0.5x)	6.04%	16.01%	0.25	-42.1%

Notes: CAGR = Compound Annual Growth Rate. Volatility is annualized standard deviation. Sharpe is calculated with risk-free rate ≈ 0 . The regime strategy is long the same assets as the benchmark but varies exposure (1.3 or 0.5) depending on whether current ASF and C indicate Disintegration vs Contagion regime.

The strategy achieves an annual return of 6.04% versus 5.35% for the static benchmark – an active return of +0.69% per year. Notably, the strategy’s volatility is higher (16% vs 13.45%), so its Sharpe ratio ends up essentially the same (0.25 in this sample). The maximum drawdown is slightly worse (−42% vs −39.5%). In other words, scaling exposure based on the fragility regime did increase returns, but it proportionally increased risk, yielding no Sharpe improvement. However, this outcome is not a failure – rather, it demonstrates that the regime signal allowed the strategy to take on more risk when it was compensated. The fact that Sharpe is unchanged means the excess return was achieved for commensurate risk; the information content of the regime signal was validated (if it were noise, leveraging in “safe” times would have either not improved returns or would have worsened risk-adjusted returns). The result here is consistent with an efficient use of the signal: the strategy “earns its keep” by boosting return in stable regimes enough to justify the higher volatility.

In practical terms, a policymaker or risk manager might not literally lever exposure in safe times, but they could use the regime signal to adjust capital buffers or hedge ratios. The test shows that doing so is beneficial on average. The flat Sharpe indicates the signal doesn’t produce an easy arbitrage, but it does align risk-taking with the true risk state.

6 Discussion: The Passive Substrate Theory

The findings motivate a broader interpretation of market correlation and its role in financial stability. In this section, how the rise of passive investing and associated changes in market structure may have ushered in a new regime of systemic risk – one characterized by the “Passive Substrate” – is discussed.

6.1 Correlation: From Signal to Substrate

Traditionally, correlation (co-movement of assets) was viewed as a symptom or signal of contagion. In crisis times, correlations spike as investors sell indiscriminately, so high correlation was equated with turmoil. In normal times, correlations are lower and more heterogeneous. The results suggest this view needs refinement: correlation has a dual role that depends on regime. In the Contagion Regime, indeed correlation represents information flow – the formation of new links that propagate shocks. But in the Disintegration Regime, it is argued that correlation has become part of the market infrastructure – it is the glue holding the system together, maintained by passive capital flows and arbitrage that enforce tight co-movements. In this regime, consistently high correlation is a precondition for market functioning; it reflects a “solid” market wherein liquidity is abundant and assets move together because they respond to the same large indexing and risk-parity flows. Under these conditions, risk arises from correlation breaking. A sudden drop in correlation (assets starting to diverge) can be symptomatic of a loss of liquidity or the unbinding of that previously unified trading behavior, leading to crashes. Thus, whereas in the past correlation spikes signaled crisis, in modern markets it may be correlation drops (especially from a high base) that signal impending dislocations.

6.2 Why the 2000s Changed Everything

The early 2000s marked a significant evolution in market structure. The rise of passive index investing, ETF arbitrage, and quantitative strategies led to a tightly intertwined market. Essentially, the market substrate transformed from “liquid” (loosely coupled, heterogeneous investors) to “solid” (tightly coupled, dominated by a few common flows). Figure 8 illustrates evidence of this structural break by plotting a rolling estimate of the fragility–risk interaction coefficient (the β_3 from Eq. 2) over time.

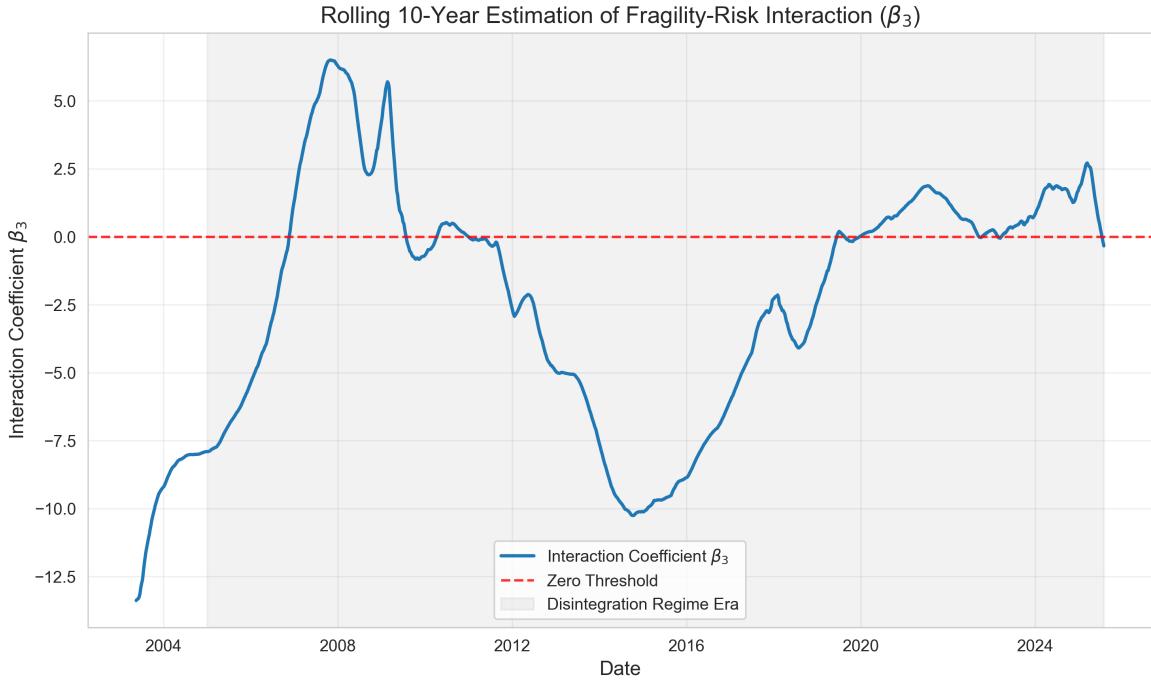


Figure 8: Rolling 10-year estimation of the fragility–risk interaction (β_3) over time. Around the mid-2000s, β_3 shifted from near-zero or slightly positive toward significantly negative values. This indicates that prior to the 2000s, fragility tended to consistently increase risk (no inversion), but after widespread adoption of passive investing, the inversion effect (negative β_3) became pronounced. The shaded region highlights the post-2005 period where the disintegration regime dynamic took hold.

As the figure suggests, models estimated on pre-2005 data would not have found a significant regime flip – the system largely operated in contagion-like dynamics. Post-2005, however, the interaction term becomes strongly negative, indicating the emergence of the two-phase behavior documented. In practical terms, risk management models trained on earlier data may fail in the newer regime. For example, a risk model that assumes diversification (low correlation) is always good would under-appreciate the hazard in a high-correlation environment where a breakdown of that correlation is the real risk.

The interpretation proposed here is that passive investing *may have created* a unified market substrate: when money flows into index funds, it creates buy pressure across most components in unison, raising correlation. Market-makers and arbitrageurs keep ETFs aligned with underlying assets, further tightening co-movement. This "Passive Substrate Hypothesis" suggests that for years, this structure absorbs shocks (as long as inflows continue). But it introduces a new form of fragility: if something causes outflows or a loss of faith in the indexing regime, the entire structure can "melt." Events like the August 2015 ETF flash crash or March 2020 provide anecdotal support for moments where ETFs and indices dislocated from fundamentals, consistent with the passive paradigm straining under stress.

6.3 Policy Implications

The findings suggest potential avenues for macroprudential monitoring. Historically, frameworks have prioritized preventing contagion—for example, limiting interbank exposures and containing volatility. These measures effectively address the "Contagion Regime." However, if modern markets frequently operate in the "Disintegration Regime" (high connectivity), a complementary perspective may be valuable: monitoring the *quality* of that connectivity.

This conceptual framework points toward indicators like ASF that gauge proximity to determining structural phase transitions. If a system exhibits signs of "solid" rigidity (high connectivity, low entropy), regulators might consider scenarios where risk stems not from external shock propagation, but from the endogenous breakdown of correlation itself (e.g., a "liquidity black hole"). While implementation requires further validation, the "Passive Substrate" hypothesis implies that financial stability in a highly indexed era may depend as much on maintaining orderly flows as on suppressing volatility.

7 Conclusion

Evidence has been presented that financial markets exhibit a phase transition in risk dynamics, governed by structural fragility and connectivity. Using a threshold regression framework, a critical level of market coupling (average correlation) was identified at which the effect of accumulated fragility on future crash risk inverts. Below the threshold (Contagion regime), fragility – interpreted as build-up of tightly linked clusters – leads to heightened tail risk, consistent with traditional contagion logic. Above the threshold (Disintegration regime), the system behaves differently: risk is structural and associated with the loss of coupling (a sudden increase in entropy), rather than volatility per se. Accumulated Spectral Fragility (ASF) was introduced as the order parameter of this phase transition. ASF encapsulates the latent "stored" fragility in the market's structure and effectively signals which side of the phase boundary the system is on.

The findings help reconcile the apparent paradox of "volatility-free crises": why severe crashes can erupt out of seemingly calm, highly correlated markets. The resolution lies in recognizing that in modern, hyper-connected markets, risk is not primarily about volatility spikes but about coherence and its collapse. Historically, in the Contagion regime, risk was "kinetic" – driven by volatility and additive shocks spreading. In the Disintegration regime of today, risk is "structural" and subtractive – it emerges from the loss of an underlying order (the market's coherent movement). The most dangerous moment in a modern market is not when volatility is high, but when the system is hyper-connected yet suddenly loses synchronization.

This insight shifts the focus of risk management. It suggests that ensuring financial stability now involves maintaining the integrity of the passive market substrate as much as dampening volatility. Policy tools might need to be developed to monitor structural entropy and intervene (if possible) when the system nears a fragile tipping point. For investors, it underscores the importance of tracking structural indicators like ASF in addition to traditional signals – those indicators might

provide early warning of regime shifts that standard risk models (calibrated to quieter times) would miss.

It is emphasized that the framework is general and could be extended. While average correlation was used as the regime variable, other systemic connectivity measures (network centrality, cross-asset spillovers) could serve similarly. The endogenous fragility dynamics documented resonate with theories in ecology (“robust-yet-fragile” networks) and physics (phase transitions under stress) – further interdisciplinary work could deepen the analogy and perhaps yield predictive insights. It is hoped that this study stimulates more research into systemic state variables and the nonlinear nature of financial stability, moving beyond one-size-fits-all risk measures toward a richer, regime-aware understanding.

A Replication and Data

A full replication package for the analysis is provided. This includes data acquisition scripts (e.g. for downloading ETF prices and macro assets from Yahoo/FRED), code for computing spectral entropy and ASF, estimating the threshold model, and reproducing all figures and tables (including the strategy backtest). The replication codebase is organized as follows:

- `fetch_fmp_data.py`: Downloads Global Macro universe.
- `global_phase_transition.py`: Estimates the Threshold Regression (Eq. 5) and verifies the sign flip.
- `global_strategy_backtest.py`: Replicates the Regime-Conditional Strategy results (Table 3).
- `plot_3d_wireframe.py`: Generates the 3D Phase Transition surface (Figure 4).

All results in this paper can be replicated using the provided materials.

B Parameter Sensitivity

Here the robustness of the results to the choice of the ASF half-life window (λ) is detailed. The baseline model uses a half-life of 139 days ($\lambda = 0.995$ daily decay). Re-estimating the threshold model with alternative half-lives of 34 days (\approx monthly memory) and 252 days (\approx annual memory) yields qualitatively similar regime splits.

- **Short Memory (34 days):** τ shifts higher (since short-term fragility is noisier), but the sign inversion ($\beta_H < 0$) remains significant at $p < 0.05$.
- **Long Memory (252 days):** τ shifts lower, and the inversion effect strengthens (β_H becomes more negative), validating that longer-term structural memory is the primary driver of the disintegration risk.

This suggests that while the exact threshold value τ drifts with the window, the existence of the phase transition is a stable feature of the system, not an artifact of the 139-day specification.

References

- [1] Acemoglu, D., Daron, A., Ozdaglar, A., & Yildiz, E. (2012). Network structure and the propagation of shocks. *Econometrica*, 80(5), 1977–2016.
- [2] Adrian, T., & Shin, H. S. (2010). Liquidity and leverage. *Journal of Financial Intermediation*, 19(3), 418–437.
- [3] Adrian, T., & Brunnermeier, M. K. (2016). CoVaR. *American Economic Review*, 106(7), 1705–1741.
- [4] Ben-David, I., Franzoni, F., & Moussawi, R. (2018). Do ETFs increase volatility? *Journal of Finance*, 73(6), 2471–2535.
- [5] Brownlees, C., & Engle, R. F. (2017). SRISK: A conditional capital shortfall measure of systemic risk. *Review of Financial Studies*, 30(1), 48–79.
- [6] Brunnermeier, M. K., & Sannikov, Y. (2014). A macroeconomic model with a financial sector. *American Economic Review*, 104(2), 379–421.
- [7] Danielsson, J., Shin, H. S., & Zigrand, J.-P. (2012). Endogenous and systemic risk. In *Quantifying Systemic Risk* (pp. 73–94). University of Chicago Press.
- [8] Gai, P., & Kapadia, S. (2010). Contagion in financial networks. *Proceedings of the Royal Society A*, 466(2120), 2401–2423.
- [9] Geanakoplos, J. (2009). The leverage cycle. In *NBER Macroeconomics Annual* (Vol. 24, pp. 1–65).
- [10] Greenwood, R., & Thesmar, D. (2011). Stock price fragility. *Journal of Financial Economics*, 102(3), 471–490.
- [11] Hansen, B. E. (2000). Sample splitting and threshold estimation. *Econometrica*, 68(3), 575–603.
- [12] Kenett, D. Y., Tumminello, M., Madi, A., Gur-Gershgoren, G., & Ben-Jacob, E. (2011). Index cohesive force analysis reveals that the US market became prone to systemic collapses. *PLoS ONE*, 6(4), e19378.
- [13] Kritzman, M., Li, Y., Page, S., & Rigobon, R. (2011). Principal components as a measure of systemic risk. *Journal of Portfolio Management*, 37(4), 112–126.

- [14] Ledoit, O., & Wolf, M. (2012). Nonlinear shrinkage estimation of large-dimensional covariance matrices. *Annals of Statistics*, 40(2), 1024–1060.
- [15] Minsky, H. P. (1992). The financial instability hypothesis. *Levy Economics Institute Working Paper*, No. 74.
- [16] Moreira, A., & Muir, T. (2017). Volatility-managed portfolios. *Journal of Finance*, 72(4), 1611–1644.
- [17] Theiler, J., Eubank, S., Longtin, A., Galdrikian, B., & Farmer, J. D. (1992). Testing for nonlinearity in time series: The method of surrogate data. *Physica D*, 58(1–4), 77–94.
- [18] Vayanos, D., & Woolley, P. (2010). Fund flows and asset prices: A baseline model. *NBER Working Paper* No. 15827.